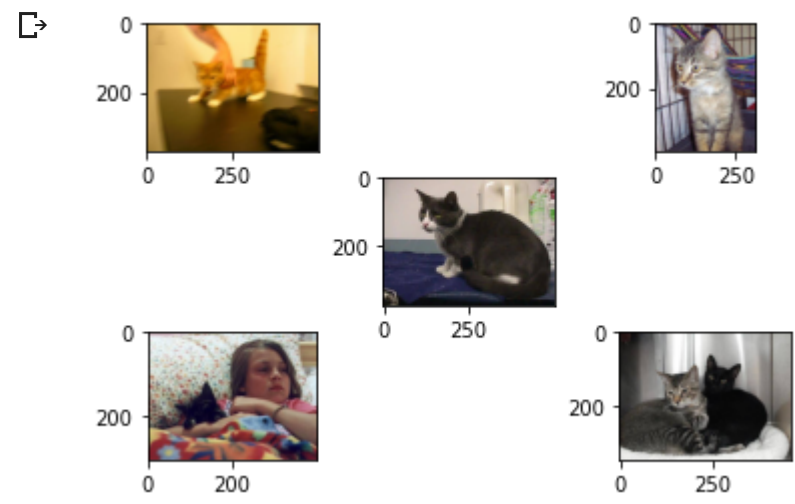




plt.show()



```
from matplotlib import pyplot
import numpy as np
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import Flatten
from keras.optimizers import SGD
import tensorflow as tf
from keras.preprocessing.image import ImageDataGenerator
```

## IMAGE AUGMENTATON

```
train_datagen = ImageDataGenerator(rescale = 1.0/255.0,
                                   shear_range = 0.2,
                                   zoom_range = 0.2,
                                   horizontal_flip = True,
                                   )

training_set = train_datagen.flow_from_directory('./dataset',
                                                target_size = (48, 48),
                                                batch_size = 64,
                                                class_mode = 'binary')

validation_datagen = ImageDataGenerator(rescale = 1.0/255.0,)
validation_set = validation_datagen.flow_from_directory('./validation',
                                                        target_size = (48, 48),
                                                        batch_size = 64,
                                                        class_mode = 'binary')
```

```
Found 20000 images belonging to 2 classes.
Found 5000 images belonging to 2 classes.
```

## MODEL 1 - OPTIMIZER : NADAM

```
model = Sequential()
model.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2,2)))

model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2,2)))

model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2,2)))

model.add(Flatten())

model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Nadam(
    learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07)
model.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model.summary()
```




```
Model: "sequential_12"
Layer (type)                Output Shape              Param #
-----
from keras.callbacks import ModelCheckpoint

checkpoint = ModelCheckpoint(filepath = 'best_model_NADAM.h5',save_best_only = True,verbose=1)

history_nadam = model.fit(training_set,
                           epochs=10,
                           callbacks=[checkpoint],
                           validation_data=validation_set)
```

Epoch 1/10  
313/313 [=====] - ETA: 0s - loss: 0.7196 - accuracy: 0.5142  
Epoch 00001: val\_loss improved from inf to 0.69263, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 84s 269ms/step - loss: 0.7196 - accuracy: 0.5142 - val\_loss: 0.6926 - val\_accuracy: 0.5000  
Epoch 2/10  
313/313 [=====] - ETA: 0s - loss: 0.6893 - accuracy: 0.5347  
Epoch 00002: val\_loss improved from 0.69263 to 0.68071, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 82s 261ms/step - loss: 0.6893 - accuracy: 0.5347 - val\_loss: 0.6807 - val\_accuracy: 0.5548  
Epoch 3/10  
313/313 [=====] - ETA: 0s - loss: 0.6702 - accuracy: 0.5985  
Epoch 00003: val\_loss improved from 0.68071 to 0.60553, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 80s 257ms/step - loss: 0.6702 - accuracy: 0.5985 - val\_loss: 0.6055 - val\_accuracy: 0.6738  
Epoch 4/10  
313/313 [=====] - ETA: 0s - loss: 0.5851 - accuracy: 0.6899  
Epoch 00004: val\_loss improved from 0.60553 to 0.54122, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 80s 254ms/step - loss: 0.5851 - accuracy: 0.6899 - val\_loss: 0.5412 - val\_accuracy: 0.7232  
Epoch 5/10  
313/313 [=====] - ETA: 0s - loss: 0.5181 - accuracy: 0.7444  
Epoch 00005: val\_loss improved from 0.54122 to 0.46390, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 80s 255ms/step - loss: 0.5181 - accuracy: 0.7444 - val\_loss: 0.4639 - val\_accuracy: 0.7826  
Epoch 6/10  
313/313 [=====] - ETA: 0s - loss: 0.4558 - accuracy: 0.7839  
Epoch 00006: val\_loss did not improve from 0.46390  
313/313 [=====] - 79s 252ms/step - loss: 0.4558 - accuracy: 0.7839 - val\_loss: 0.6407 - val\_accuracy: 0.6882  
Epoch 7/10  
313/313 [=====] - ETA: 0s - loss: 0.4069 - accuracy: 0.8137  
Epoch 00007: val\_loss improved from 0.46390 to 0.38349, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 79s 251ms/step - loss: 0.4069 - accuracy: 0.8137 - val\_loss: 0.3835 - val\_accuracy: 0.8256  
Epoch 8/10  
313/313 [=====] - ETA: 0s - loss: 0.3555 - accuracy: 0.8403  
Epoch 00008: val\_loss improved from 0.38349 to 0.34450, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 78s 250ms/step - loss: 0.3555 - accuracy: 0.8403 - val\_loss: 0.3445 - val\_accuracy: 0.8512  
Epoch 9/10  
313/313 [=====] - ETA: 0s - loss: 0.3301 - accuracy: 0.8531  
Epoch 00009: val\_loss improved from 0.34450 to 0.29266, saving model to best\_model\_NADAM.h5  
313/313 [=====] - 78s 251ms/step - loss: 0.3301 - accuracy: 0.8531 - val\_loss: 0.2927 - val\_accuracy: 0.8724  
Epoch 10/10  
313/313 [=====] - ETA: 0s - loss: 0.2985 - accuracy: 0.8706  
Epoch 00010: val\_loss did not improve from 0.29266  
313/313 [=====] - 78s 250ms/step - loss: 0.2985 - accuracy: 0.8706 - val\_loss: 0.2938 - val\_accuracy: 0.8720

```
plt.plot(history_nadam.history['accuracy'])
plt.plot(history_nadam.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model.save_weights('Atharva_Nadam_weights_48.h5')
np.save('Atharva_Nadam_48.npy',history_nadam.history)
```

## MODEL 2 - OPTIMIZER : SGD

```
model_sgd = Sequential()
model_sgd.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_sgd.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgd.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgd.add(Flatten())

model_sgd.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_sgd.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.SGD(
    learning_rate=0.01, momentum=0.0, nesterov=False)
model_sgd.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_sgd.summary()
```



Model: "sequential\_13"

Layer (type)	Output Shape	Param #
=====		
conv2d_104 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_105 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_39 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_106 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_107 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_40 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_108 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_109 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_110 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_111 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_41 (MaxPooling)	(None, 6, 6, 64)	0


from keras.callbacks import ModelCheckpoint

checkpoint = ModelCheckpoint(filepath = 'best\_model\_SGD.h5',save\_best\_only = True,verbose=1)

history\_sgd = model\_sgd.fit(training\_set,  
epochs=10,  
callbacks=[checkpoint],  
validation\_data=validation\_set)

🔗 Epoch 1/10  
313/313 [=====] - ETA: 0s - loss: 0.6796 - accuracy: 0.5662  
Epoch 00001: val\_loss improved from inf to 0.68032, saving model to best\_model\_SGD.h5  
313/313 [=====] - 79s 251ms/step - loss: 0.6796 - accuracy: 0.5662 - val\_loss: 0.6803 - val\_accuracy: 0.5606  
Epoch 2/10  
313/313 [=====] - ETA: 0s - loss: 0.6421 - accuracy: 0.6255  
Epoch 00002: val\_loss improved from 0.68032 to 0.61192, saving model to best\_model\_SGD.h5  
313/313 [=====] - 79s 251ms/step - loss: 0.6421 - accuracy: 0.6255 - val\_loss: 0.6119 - val\_accuracy: 0.6742  
Epoch 3/10  
313/313 [=====] - ETA: 0s - loss: 0.6091 - accuracy: 0.6681  
Epoch 00003: val\_loss did not improve from 0.61192  
313/313 [=====] - 79s 252ms/step - loss: 0.6091 - accuracy: 0.6681 - val\_loss: 0.6181 - val\_accuracy: 0.6420  
Epoch 4/10  
313/313 [=====] - ETA: 0s - loss: 0.5835 - accuracy: 0.6888  
Epoch 00004: val\_loss improved from 0.61192 to 0.53752, saving model to best\_model\_SGD.h5  
313/313 [=====] - 78s 250ms/step - loss: 0.5835 - accuracy: 0.6888 - val\_loss: 0.5375 - val\_accuracy: 0.7332  
Epoch 5/10  
313/313 [=====] - ETA: 0s - loss: 0.5612 - accuracy: 0.7095  
Epoch 00005: val\_loss improved from 0.53752 to 0.51241, saving model to best\_model\_SGD.h5  
313/313 [=====] - 78s 249ms/step - loss: 0.5612 - accuracy: 0.7095 - val\_loss: 0.5124 - val\_accuracy: 0.7466  
Epoch 6/10  
313/313 [=====] - ETA: 0s - loss: 0.5478 - accuracy: 0.7203  
Epoch 00006: val\_loss did not improve from 0.51241  
313/313 [=====] - 78s 249ms/step - loss: 0.5478 - accuracy: 0.7203 - val\_loss: 0.5565 - val\_accuracy: 0.7140  
Epoch 7/10  
313/313 [=====] - ETA: 0s - loss: 0.5281 - accuracy: 0.7372  
Epoch 00007: val\_loss did not improve from 0.51241  
313/313 [=====] - 78s 249ms/step - loss: 0.5281 - accuracy: 0.7372 - val\_loss: 0.5375 - val\_accuracy: 0.7258  
Epoch 8/10  
313/313 [=====] - ETA: 0s - loss: 0.5158 - accuracy: 0.7461  
Epoch 00008: val\_loss did not improve from 0.51241  
313/313 [=====] - 78s 250ms/step - loss: 0.5158 - accuracy: 0.7461 - val\_loss: 0.6087 - val\_accuracy: 0.6914  
Epoch 9/10  
313/313 [=====] - ETA: 0s - loss: 0.4977 - accuracy: 0.7572  
Epoch 00009: val\_loss improved from 0.51241 to 0.50194, saving model to best\_model\_SGD.h5  
313/313 [=====] - 78s 250ms/step - loss: 0.4977 - accuracy: 0.7572 - val\_loss: 0.5019 - val\_accuracy: 0.7522  
Epoch 10/10  
313/313 [=====] - ETA: 0s - loss: 0.4817 - accuracy: 0.7700  
Epoch 00010: val\_loss improved from 0.50194 to 0.47040, saving model to best\_model\_SGD.h5  
313/313 [=====] - 78s 249ms/step - loss: 0.4817 - accuracy: 0.7700 - val\_loss: 0.4704 - val\_accuracy: 0.7754

plt.plot(history\_sgd.history['accuracy'])  
plt.plot(history\_sgd.history['val\_accuracy'])  
plt.title('model accuracy')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'test'], loc='upper left')  
plt.show()



model\_sgd.save\_weights('Atharva\_SGD\_48.h5')  
np.save('Atharva\_SGD\_48.npy',history\_sgd.history)

## ➤ MODEL 3 - OPTIMIZER : SGD + NESTROV

```
model_sgdn = Sequential()  
model_sgdn.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))  
model_sgdn.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_sgdn.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))  
model_sgdn.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_sgdn.add(Flatten())
```



```
model_sgdn.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
```

```
model_sgdn.add(Dense(1, activation='sigmoid'))
```

```
opt = tf.keras.optimizers.SGD(
    learning_rate=0.01, momentum=0.0, nesterov=True)
model_sgdn.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```

```
model_sgdn.summary()
```

Model: "sequential\_14"

Layer (type)	Output Shape	Param #
=====		
conv2d_112 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_113 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_42 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_114 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_115 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_43 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_116 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_117 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_118 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_119 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_44 (MaxPooling)	(None, 6, 6, 64)	0
flatten_14 (Flatten)	(None, 2304)	0
dense_28 (Dense)	(None, 128)	295040
dense_29 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		
=====		


```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath = 'best_model_SGDN.h5',save_best_only = True,verbose=1)
```

```
history_sgdn = model_sgdn.fit(training_set,
                               epochs=10,
                               callbacks=[checkpoint],
                               validation_data=validation_set)
```

Epoch 1/10  
313/313 [=====] - ETA: 0s - loss: 0.6764 - accuracy: 0.5730  
Epoch 00001: val\_loss improved from inf to 0.65042, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 79s 252ms/step - loss: 0.6764 - accuracy: 0.5730 - val\_loss: 0.6504 - val\_accuracy: 0.6062  
Epoch 2/10  
313/313 [=====] - ETA: 0s - loss: 0.6369 - accuracy: 0.6341  
Epoch 00002: val\_loss improved from 0.65042 to 0.61374, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 78s 250ms/step - loss: 0.6369 - accuracy: 0.6341 - val\_loss: 0.6137 - val\_accuracy: 0.6496  
Epoch 3/10  
313/313 [=====] - ETA: 0s - loss: 0.6067 - accuracy: 0.6701  
Epoch 00003: val\_loss did not improve from 0.61374  
313/313 [=====] - 78s 250ms/step - loss: 0.6067 - accuracy: 0.6701 - val\_loss: 0.7182 - val\_accuracy: 0.6076  
Epoch 4/10  
313/313 [=====] - ETA: 0s - loss: 0.5873 - accuracy: 0.6909  
Epoch 00004: val\_loss improved from 0.61374 to 0.55898, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 78s 250ms/step - loss: 0.5873 - accuracy: 0.6909 - val\_loss: 0.5590 - val\_accuracy: 0.7130  
Epoch 5/10  
313/313 [=====] - ETA: 0s - loss: 0.5642 - accuracy: 0.7089  
Epoch 00005: val\_loss improved from 0.55898 to 0.54637, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 78s 248ms/step - loss: 0.5642 - accuracy: 0.7089 - val\_loss: 0.5464 - val\_accuracy: 0.7266  
Epoch 6/10  
313/313 [=====] - ETA: 0s - loss: 0.5468 - accuracy: 0.7196  
Epoch 00006: val\_loss improved from 0.54637 to 0.53545, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 77s 246ms/step - loss: 0.5468 - accuracy: 0.7196 - val\_loss: 0.5355 - val\_accuracy: 0.7362  
Epoch 7/10  
313/313 [=====] - ETA: 0s - loss: 0.5309 - accuracy: 0.7347  
Epoch 00007: val\_loss improved from 0.53545 to 0.49420, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 77s 247ms/step - loss: 0.5309 - accuracy: 0.7347 - val\_loss: 0.4942 - val\_accuracy: 0.7568  
Epoch 8/10  
313/313 [=====] - ETA: 0s - loss: 0.5162 - accuracy: 0.7445  
Epoch 00008: val\_loss improved from 0.49420 to 0.47550, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 77s 248ms/step - loss: 0.5162 - accuracy: 0.7445 - val\_loss: 0.4755 - val\_accuracy: 0.7768  
Epoch 9/10  
313/313 [=====] - ETA: 0s - loss: 0.4961 - accuracy: 0.7582  
Epoch 00009: val\_loss improved from 0.47550 to 0.47301, saving model to best\_model\_SGDN.h5  
313/313 [=====] - 77s 247ms/step - loss: 0.4961 - accuracy: 0.7582 - val\_loss: 0.4730 - val\_accuracy: 0.7686  
Epoch 10/10  
313/313 [=====] - ETA: 0s - loss: 0.4852 - accuracy: 0.7679  
Epoch 00010: val\_loss did not improve from 0.47301  
313/313 [=====] - 77s 246ms/step - loss: 0.4852 - accuracy: 0.7679 - val\_loss: 0.5043 - val\_accuracy: 0.7514

```
plt.plot(history_sgdn.history['accuracy'])
plt.plot(history_sgdn.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_sgdn.save_weights('Atharva_SGDN_48.h5')
np.save('Atharva_SGDN_48.npy',history_sgdn.history)
```

MODEL 4 - OPTIMIZER : RMS

```
model_rms = Sequential()
model_rms.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_rms.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(MaxPooling2D((2, 2), strides=(2,2)))

model_rms.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(MaxPooling2D((2, 2), strides=(2,2)))

model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(MaxPooling2D((2, 2), strides=(2,2)))

model_rms.add(Flatten())

model_rms.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_rms.add(Dense(1, activation='sigmoid'))
```

```
opt = tf.keras.optimizers.RMSprop(
    learning_rate=0.001,
    rho=0.9,
    momentum=0.0,
    epsilon=1e-07,
    centered=False,)
model_rms.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_rms.summary()
```

Model: "sequential\_15"

Layer (type)	Output Shape	Param #
conv2d_120 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_121 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_45 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_122 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_123 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_46 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_124 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_125 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_126 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_127 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_47 (MaxPooling)	(None, 6, 6, 64)	0
flatten_15 (Flatten)	(None, 2304)	0
dense_30 (Dense)	(None, 128)	295040
dense_31 (Dense)	(None, 1)	129
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		

```
from keras.callbacks import ModelCheckpoint

checkpoint = ModelCheckpoint(filepath = 'best_model_RMS.h5',save_best_only = True,verbose=1)

history_rms = model_rms.fit(training_set,
                             epochs=10,
                             callbacks=[checkpoint],
                             validation_data=validation_set)
```

Epoch 1/10

313/313 [=====] - ETA: 0s - loss: 1.1355 - accuracy: 0.5162

Epoch 00001: val\_loss improved from inf to 0.69124, saving model to best\_model\_RMS.h5

313/313 [=====] - 78s 248ms/step - loss: 1.1355 - accuracy: 0.5162 - val\_loss: 0.6912 - val\_accuracy: 0.5010

Epoch 2/10

313/313 [=====] - ETA: 0s - loss: 0.6784 - accuracy: 0.5974

Epoch 00002: val\_loss improved from 0.69124 to 0.59583, saving model to best\_model\_RMS.h5

313/313 [=====] - 77s 246ms/step - loss: 0.6784 - accuracy: 0.5974 - val\_loss: 0.5958 - val\_accuracy: 0.6910

Epoch 3/10

313/313 [=====] - ETA: 0s - loss: 0.5916 - accuracy: 0.6934

Epoch 00003: val\_loss improved from 0.59583 to 0.51424, saving model to best\_model\_RMS.h5

313/313 [=====] - 78s 248ms/step - loss: 0.5916 - accuracy: 0.6934 - val\_loss: 0.5142 - val\_accuracy: 0.7472

Epoch 4/10

313/313 [=====] - ETA: 0s - loss: 0.5089 - accuracy: 0.7548

Epoch 00004: val\_loss improved from 0.51424 to 0.47275, saving model to best\_model\_RMS.h5

313/313 [=====] - 78s 248ms/step - loss: 0.5089 - accuracy: 0.7548 - val\_loss: 0.4728 - val\_accuracy: 0.7954

Epoch 5/10

313/313 [=====] - ETA: 0s - loss: 0.4463 - accuracy: 0.7972

Epoch 00005: val\_loss improved from 0.47275 to 0.37714, saving model to best\_model\_RMS.h5

313/313 [=====] - 77s 246ms/step - loss: 0.4463 - accuracy: 0.7972 - val\_loss: 0.3771 - val\_accuracy: 0.8228

Epoch 6/10

313/313 [=====] - ETA: 0s - loss: 0.4009 - accuracy: 0.8160

Epoch 00006: val\_loss improved from 0.37714 to 0.34803, saving model to best\_model\_RMS.h5

313/313 [=====] - 77s 246ms/step - loss: 0.4009 - accuracy: 0.8160 - val\_loss: 0.3480 - val\_accuracy: 0.8480

Epoch 7/10

313/313 [=====] - ETA: 0s - loss: 0.3708 - accuracy: 0.8323

Epoch 00007: val\_loss improved from 0.34803 to 0.31253, saving model to best\_model\_RMS.h5

313/313 [=====] - 77s 246ms/step - loss: 0.3708 - accuracy: 0.8323 - val\_loss: 0.3125 - val\_accuracy: 0.8616

Epoch 8/10

313/313 [=====] - ETA: 0s - loss: 0.3437 - accuracy: 0.8490

Epoch 00008: val\_loss improved from 0.31253 to 0.31074, saving model to best\_model\_RMS.h5

313/313 [=====] - 77s 245ms/step - loss: 0.3437 - accuracy: 0.8490 - val\_loss: 0.3107 - val\_accuracy: 0.8694

Epoch 9/10

313/313 [=====] - ETA: 0s - loss: 0.3290 - accuracy: 0.8562

Epoch 00009: val\_loss did not improve from 0.31074

313/313 [=====] - 77s 245ms/step - loss: 0.3290 - accuracy: 0.8562 - val\_loss: 0.3168 - val\_accuracy: 0.8764

Epoch 10/10


313/313 [=====] - ETA: 0s - loss: 0.3126 - accuracy: 0.8663

Epoch 00010: val\_loss improved from 0.31074 to 0.28536, saving model to best\_model\_RMS.h5

313/313 [=====] - 77s 245ms/step - loss: 0.3126 - accuracy: 0.8663 - val\_loss: 0.2854 - val\_accuracy: 0.8818

```
plt.plot(history_rms.history['accuracy'])
plt.plot(history_rms.history['val_accuracy'])
plt.title('model accuracy')
```

```
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_rms.save_weights('Atharva_RMS_4850.h5')
np.save('Atharva_RMS_4850.npy',history_rms.history)
```

MODEL 5 - OPTIMIZER : ADAM

```
model_adam = Sequential()
model_adam.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adam.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adam.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adam.add(Flatten())

model_adam.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adam.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Adam(
    learning_rate=0.001,
    beta_1=0.9,
    beta_2=0.999,
    epsilon=1e-07,
    amsgrad=False,)
model_adam.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_adam.summary()
```

Model: "sequential\_16"

Layer (type)	Output Shape	Param #
=====		
conv2d_128 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_129 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_48 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_130 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_131 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_49 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_132 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_133 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_134 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_135 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_50 (MaxPooling)	(None, 6, 6, 64)	0
flatten_16 (Flatten)	(None, 2304)	0
dense_32 (Dense)	(None, 128)	295040
dense_33 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		

```
from keras.callbacks import ModelCheckpoint


checkpoint = ModelCheckpoint(filepath = 'best_model_ADAM.h5',save_best_only = True,verbose=1)

history_adam = model_adam.fit(training_set,
    epochs=10,
    callbacks=[checkpoint],
    validation_data=validation_set)
```



```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6936 - accuracy: 0.5601
Epoch 00001: val_loss improved from inf to 0.65515, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 247ms/step - loss: 0.6936 - accuracy: 0.5601 - val_loss: 0.6551 - val_accuracy: 0.6302
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6354 - accuracy: 0.6440
Epoch 00002: val_loss improved from 0.65515 to 0.58563, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.6354 - accuracy: 0.6440 - val_loss: 0.5856 - val_accuracy: 0.6976
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.5680 - accuracy: 0.7085
Epoch 00003: val_loss improved from 0.58563 to 0.50897, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.5680 - accuracy: 0.7085 - val_loss: 0.5090 - val_accuracy: 0.7546
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.4974 - accuracy: 0.7588
Epoch 00004: val_loss improved from 0.50897 to 0.44578, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.4974 - accuracy: 0.7588 - val_loss: 0.4458 - val_accuracy: 0.7978
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.4427 - accuracy: 0.7914
Epoch 00005: val_loss improved from 0.44578 to 0.39582, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 246ms/step - loss: 0.4427 - accuracy: 0.7914 - val_loss: 0.3958 - val_accuracy: 0.8172
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.3863 - accuracy: 0.8246
Epoch 00006: val_loss improved from 0.39582 to 0.34291, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.3863 - accuracv: 0.8246 - val loss: 0.3429 - val accuracv: 0.8506
```

```
plt.plot(history_adam.history['accuracy'])
plt.plot(history_adam.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adam.save_weights('Atharva_ADAM_48.h5')
np.save('Atharva_ADAM_48.npy',history_adam.history)
```

## MODEL 6 - OPTIMIZER : ADADelta

```
model_adad = Sequential()
model_adad.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adad.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adad.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adad.add(Flatten())

model_adad.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adad.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Adadelta(
    learning_rate=0.001, rho=0.95, epsilon=1e-07, name="Adadelta", )
model_adad.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_adad.summary()
```

Model: "sequential_17"		
Layer (type)	Output Shape	Param #
=====		
conv2d_136 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_137 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_51 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_138 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_139 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_52 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_140 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_141 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_142 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_143 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_53 (MaxPooling)	(None, 6, 6, 64)	0
flatten_17 (Flatten)	(None, 2304)	0
dense_34 (Dense)	(None, 128)	295040
dense_35 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		

```
from keras.callbacks import ModelCheckpoint
```


```
checkpoint = ModelCheckpoint(filepath = 'best_model_ADADELTA.h5',save_best_only = True,verbose=1)
```



```
history_adad = model_adad.fit(training_set,
                               epochs=10,
                               callbacks=[checkpoint],
                               validation_data=validation_set)
```

```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6824 - accuracy: 0.5621
Epoch 00001: val_loss improved from inf to 0.67433, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 251ms/step - loss: 0.6824 - accuracy: 0.5621 - val_loss: 0.6743 - val_accuracy: 0.5886
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6702 - accuracy: 0.6010
Epoch 00002: val_loss improved from 0.67433 to 0.66563, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6702 - accuracy: 0.6010 - val_loss: 0.6656 - val_accuracy: 0.5978
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.6615 - accuracy: 0.6190
Epoch 00003: val_loss improved from 0.66563 to 0.65408, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6615 - accuracy: 0.6190 - val_loss: 0.6541 - val_accuracy: 0.6338
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.6528 - accuracy: 0.6294
Epoch 00004: val_loss improved from 0.65408 to 0.65017, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6528 - accuracy: 0.6294 - val_loss: 0.6502 - val_accuracy: 0.6212
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.6444 - accuracy: 0.6418
Epoch 00005: val_loss improved from 0.65017 to 0.64251, saving model to best_model_ADADELTA.h5
313/313 [=====] - 80s 255ms/step - loss: 0.6444 - accuracy: 0.6418 - val_loss: 0.6425 - val_accuracy: 0.6318
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.6373 - accuracy: 0.6474
Epoch 00006: val_loss improved from 0.64251 to 0.63577, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6373 - accuracy: 0.6474 - val_loss: 0.6358 - val_accuracy: 0.6380
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.6325 - accuracy: 0.6461
Epoch 00007: val_loss improved from 0.63577 to 0.63574, saving model to best_model_ADADELTA.h5
313/313 [=====] - 80s 254ms/step - loss: 0.6325 - accuracy: 0.6461 - val_loss: 0.6357 - val_accuracy: 0.6310
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.6269 - accuracy: 0.6562
Epoch 00008: val_loss improved from 0.63574 to 0.63433, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6269 - accuracy: 0.6562 - val_loss: 0.6343 - val_accuracy: 0.6312
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.6213 - accuracy: 0.6608
Epoch 00009: val_loss did not improve from 0.63433
313/313 [=====] - 79s 253ms/step - loss: 0.6213 - accuracy: 0.6608 - val_loss: 0.6356 - val_accuracy: 0.6286
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.6180 - accuracy: 0.6644
Epoch 00010: val_loss improved from 0.63433 to 0.61984, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6180 - accuracy: 0.6644 - val_loss: 0.6198 - val_accuracy: 0.6530
```

```
plt.plot(history_adad.history['accuracy'])
plt.plot(history_adad.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adad.save_weights('Atharva_ADADELTA_48.h5')
np.save('Atharva_ADADELTA_48.npy',history_adad.history)
```

## MODEL 7 - OPTIMIZER : ADAGrad

```
model_adag = Sequential()
model_adag.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adag.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_adag.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_adag.add(Flatten())
```

```
model_adag.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
```

```
model_adag.add(Dense(1, activation='sigmoid'))
```

```
opt = tf.keras.optimizers.Adagrad(
    learning_rate=0.001,
    initial_accumulator_value=0.1,
    epsilon=1e-07,)
model_adag.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```

```
model_adag.summary()
```

Model: "sequential\_18"

Layer (type)	Output Shape	Param #
=====		
conv2d_144 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_145 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_54 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_146 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_147 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_55 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_148 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_149 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_150 (Conv2D)	(None, 12, 12, 64)	36928


```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath = 'best_model_ADAGRAD.h5',save_best_only = True,verbose=1)
```

```
history_adag = model_adag.fit(training_set,
                              epochs=10,
                              callbacks=[checkpoint],
                              validation_data=validation_set)
```

```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6747 - accuracy: 0.5795
Epoch 00001: val_loss improved from inf to 0.64739, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 78s 250ms/step - loss: 0.6747 - accuracy: 0.5795 - val_loss: 0.6474 - val_accuracy: 0.6470
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6404 - accuracy: 0.6321
Epoch 00002: val_loss did not improve from 0.64739
313/313 [=====] - 78s 250ms/step - loss: 0.6404 - accuracy: 0.6321 - val_loss: 0.7419 - val_accuracy: 0.5384
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.6191 - accuracy: 0.6536
Epoch 00003: val_loss improved from 0.64739 to 0.58838, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 78s 248ms/step - loss: 0.6191 - accuracy: 0.6536 - val_loss: 0.5884 - val_accuracy: 0.6896
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.6012 - accuracy: 0.6745
Epoch 00004: val_loss did not improve from 0.58838
313/313 [=====] - 78s 250ms/step - loss: 0.6012 - accuracy: 0.6745 - val_loss: 0.6177 - val_accuracy: 0.6554
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.5848 - accuracy: 0.6920
Epoch 00005: val_loss improved from 0.58838 to 0.57214, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 79s 252ms/step - loss: 0.5848 - accuracy: 0.6920 - val_loss: 0.5721 - val_accuracy: 0.6976
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.5678 - accuracy: 0.7046
Epoch 00006: val_loss improved from 0.57214 to 0.55453, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 79s 252ms/step - loss: 0.5678 - accuracy: 0.7046 - val_loss: 0.5545 - val_accuracy: 0.7148
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.5585 - accuracy: 0.7122
Epoch 00007: val_loss improved from 0.55453 to 0.55071, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 79s 254ms/step - loss: 0.5585 - accuracy: 0.7122 - val_loss: 0.5507 - val_accuracy: 0.7162
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.5480 - accuracy: 0.7224
Epoch 00008: val_loss improved from 0.55071 to 0.52879, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 80s 254ms/step - loss: 0.5480 - accuracy: 0.7224 - val_loss: 0.5288 - val_accuracy: 0.7358
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.5385 - accuracy: 0.7278
Epoch 00009: val_loss did not improve from 0.52879
313/313 [=====] - 79s 251ms/step - loss: 0.5385 - accuracy: 0.7278 - val_loss: 0.5564 - val_accuracy: 0.7096
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.5313 - accuracy: 0.7317
Epoch 00010: val_loss improved from 0.52879 to 0.51597, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 79s 251ms/step - loss: 0.5313 - accuracy: 0.7317 - val_loss: 0.5160 - val_accuracy: 0.7448
```

```
plt.plot(history_adag.history['accuracy'])
plt.plot(history_adag.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adag.save_weights('Atharva_ADAG_48.h5')
np.save('Atharva_ADAG_48.npy',history_adag.history)
```

## MODEL 8 - OPTIMIZER : ADAMax

```
model_adamax = Sequential()
model_adamax.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adamax.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adamax.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adamax.add(Flatten())

model_adamax.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adamax.add(Dense(1, activation='sigmoid'))
```

```
opt = tf.keras.optimizers.Adamax(
    learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07,)
model_adamax.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```

```
model_adamax.summary()
```

Model: "sequential\_19"

Layer (type)	Output Shape	Param #
=====		
conv2d_152 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_153 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_57 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_154 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_155 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_58 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_156 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_157 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_158 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_159 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_59 (MaxPooling)	(None, 6, 6, 64)	0
flatten_19 (Flatten)	(None, 2304)	0
dense_38 (Dense)	(None, 128)	295040
dense_39 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		


```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath = 'best_model_ADAMAX.h5',save_best_only = True,verbose=1)
```

```
history_adamax = model_adamax.fit(training_set,
    epochs=10,
    callbacks=[checkpoint],
    validation_data=validation_set)
```

Epoch 1/10  
313/313 [=====] - ETA: 0s - loss: 0.6777 - accuracy: 0.5832  
Epoch 00001: val\_loss improved from inf to 0.60260, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 79s 251ms/step - loss: 0.6777 - accuracy: 0.5832 - val\_loss: 0.6026 - val\_accuracy: 0.6804  
Epoch 2/10  
313/313 [=====] - ETA: 0s - loss: 0.5951 - accuracy: 0.6816  
Epoch 00002: val\_loss improved from 0.60260 to 0.56081, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 79s 252ms/step - loss: 0.5951 - accuracy: 0.6816 - val\_loss: 0.5608 - val\_accuracy: 0.7026  
Epoch 3/10  
313/313 [=====] - ETA: 0s - loss: 0.5446 - accuracy: 0.7276  
Epoch 00003: val\_loss improved from 0.56081 to 0.50568, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 80s 254ms/step - loss: 0.5446 - accuracy: 0.7276 - val\_loss: 0.5057 - val\_accuracy: 0.7524  
Epoch 4/10  
313/313 [=====] - ETA: 0s - loss: 0.5038 - accuracy: 0.7513  
Epoch 00004: val\_loss improved from 0.50568 to 0.48203, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 80s 254ms/step - loss: 0.5038 - accuracy: 0.7513 - val\_loss: 0.4820 - val\_accuracy: 0.7670  
Epoch 5/10  
313/313 [=====] - ETA: 0s - loss: 0.4622 - accuracy: 0.7778  
Epoch 00005: val\_loss improved from 0.48203 to 0.41399, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 79s 252ms/step - loss: 0.4622 - accuracy: 0.7778 - val\_loss: 0.4140 - val\_accuracy: 0.8086  
Epoch 6/10  
313/313 [=====] - ETA: 0s - loss: 0.4156 - accuracy: 0.8101  
Epoch 00006: val\_loss improved from 0.41399 to 0.41311, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 79s 252ms/step - loss: 0.4156 - accuracy: 0.8101 - val\_loss: 0.4131 - val\_accuracy: 0.8182  
Epoch 7/10  
313/313 [=====] - ETA: 0s - loss: 0.3897 - accuracy: 0.8232  
Epoch 00007: val\_loss improved from 0.41311 to 0.38135, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 78s 250ms/step - loss: 0.3897 - accuracy: 0.8232 - val\_loss: 0.3813 - val\_accuracy: 0.8284  
Epoch 8/10  
313/313 [=====] - ETA: 0s - loss: 0.3548 - accuracy: 0.8444  
Epoch 00008: val\_loss improved from 0.38135 to 0.32681, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 77s 247ms/step - loss: 0.3548 - accuracy: 0.8444 - val\_loss: 0.3268 - val\_accuracy: 0.8584  
Epoch 9/10  
313/313 [=====] - ETA: 0s - loss: 0.3409 - accuracy: 0.8500  
Epoch 00009: val\_loss improved from 0.32681 to 0.31083, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 77s 246ms/step - loss: 0.3409 - accuracy: 0.8500 - val\_loss: 0.3108 - val\_accuracy: 0.8640  
Epoch 10/10  
313/313 [=====] - ETA: 0s - loss: 0.3217 - accuracy: 0.8589  
Epoch 00010: val\_loss improved from 0.31083 to 0.28348, saving model to best\_model\_ADAMAX.h5  
313/313 [=====] - 77s 246ms/step - loss: 0.3217 - accuracy: 0.8589 - val\_loss: 0.2835 - val\_accuracy: 0.8754

```
plt.plot(history_adamax.history['accuracy'])
plt.plot(history_adamax.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adamax.save_weights('Atharva_ADAMAX_48.h5')
np.save('Atharva_ADAMAX_48.npy',history_adamax.history)
```

## Testing the Model against custom test data

We will validate the test accuracy using the model generated using the RMSProp optimizer as it has the highest degree of performance

```
import cv2
import glob
predict_files = glob.glob("/content/test/*.jpg")
predictor, image_id = [], []
for file in predict_files:
    img = cv2.imread(file)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img,(48,48))
    img = np.expand_dims(img, 0)
    outcome = [np.argmax(model_rms.predict(img))]
    predictor.extend(list(outcome))
    image_id.extend([i.rsplit("\\")[-1]])
```

```
predict_files = glob.glob("/content/validation/test/*.jpg*")
predictor, image_id = [], []
for file in predict_files:
    image = cv2.imread(file)
    img = cv2.resize(image, (48,48))
    img = np.expand_dims(img, 0)

    prediction = model_rms.predict(img)

    predict = str(model_rms.predict_classes(img)[0][0])

    predictor.extend(list(predict))
```

```
predictions =[]
c = 0
w = 0
classified = ['cat','dog']
for i in range(len(predictor)):
    if predictor[i] == '0':
        predictions.append(classified[0])
    else:
        predictions.append(classified[1])
```


```
data = []
for file in predict_files:

    data.append(os.path.basename(file)[:3])

df = pd.DataFrame(data, columns=['FileName'])

df['Predictions'] = predictions

df.head(15)
```



	FileName	Predictions
0	cat	cat
1	cat	cat
2	cat	cat
3	cat	cat
4	cat	cat
5	cat	dog
6	cat	cat
7	cat	cat
8	cat	cat
9	cat	cat
10	cat	cat
11	cat	cat
12	cat	cat
13	cat	cat
14	cat	cat