

Name - Atharva Gondkar

Class - IS2

Roll Number - 2176032

Enrollment Number - MITU17BTMA0013

EXPERIMENT #3

DATE: 11-09-2020

TITLE : ML II ASSIGNMENT 3

AIM

Perform Image Classification using CatsVsDogs dataset.

OBJECTIVE

- 1. Implement CNN for Image Classification using CatsVsDogs dataset.

DRIVE LINK - <https://drive.google.com/drive/u/0/folders/11s0LnDfN-u8BMit5F-ZcLMOLsEQNV22L>

**Notebook, code, pdf, output snapshots have been stored on the above given drive link.*

```
!pip3 install patool
!pip3 install pyunpack
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
from pyunpack import Archive
Archive("/content/gdrive/My Drive/ML_LAB_ASSIGNMENT_2020/ASSIGNMENT_3/CatvsDogs.rar").extractall(".")
```

```
!mkdir dataset
!mkdir ./dataset/cat
!mkdir ./dataset/dog
!mkdir validation
!mkdir ./validation/cat
!mkdir ./validation/dog
```

```
from os import makedirs
from os import listdir
from shutil import copyfile
from random import seed
from random import random
```

```
train_src_directory = './train'
cnt_cat = 0
cnt_dog = 0
for file in listdir(train_src_directory):
    src = train_src_directory + '/' + file
    if file.startswith('cat.'):
        if cnt_cat < 10000:
            dst = 'dataset/cat/'+ file
        else:
            dst = 'validation/cat/'+ file
        copyfile(src, dst)
        cnt_cat+=1

    elif file.startswith('dog.'):
        if cnt_dog < 10000:
            dst = 'dataset/dog/'+ file
        else:
            dst = 'validation/dog/'+ file
        copyfile(src, dst)
        cnt_dog+=1
```

```
from matplotlib import pyplot as plt
from matplotlib.image import imread
```

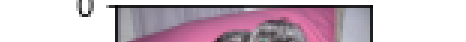


IMAGE AUGMENTATION

[illegible]

```
validation_datagen = ImageDataGenerator(rescale = 1.0/255.0,)
validation_set = validation_datagen.flow_from_directory('./validation',
                                                         target_size = (48, 48),
                                                         batch_size = 64,
                                                         class_mode = 'binary')
```

Found 20000 images belonging to 2 classes.
Found 5000 images belonging to 2 classes.

MODEL 1 - OPTIMIZER : NADAM

```
model = Sequential()
model.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2,2)))

model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2,2)))

model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model.add(MaxPooling2D((2, 2), strides=(2,2)))

model.add(Flatten())

model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Nadam(
    learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07)
model.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 48, 48, 256)	7168
conv2d_1 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d (MaxPooling2D)	(None, 24, 24, 256)	0
conv2d_2 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_3 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 128)	0
conv2d_4 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_5 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_6 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_7 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 128)	295040
dense_1 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		
=====		


```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath = 'best_model_NADAM.h5',save_best_only = True,verbose=1)
```

```
history_nadam = model.fit(training_set,
                           epochs=10,
                           callbacks=[checkpoint],
                           validation_data=validation_set)
```

```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.7196 - accuracy: 0.5142
Epoch 00001: val_loss improved from inf to 0.69263, saving model to best_model_NADAM.h5
313/313 [=====] - 84s 269ms/step - loss: 0.7196 - accuracy: 0.5142 - val_loss: 0.6926 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6893 - accuracy: 0.5347
Epoch 00002: val_loss improved from 0.69263 to 0.68071, saving model to best_model_NADAM.h5
313/313 [=====] - 82s 261ms/step - loss: 0.6893 - accuracy: 0.5347 - val_loss: 0.6807 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.6702 - accuracy: 0.5985
Epoch 00003: val_loss improved from 0.68071 to 0.60553, saving model to best_model_NADAM.h5
313/313 [=====] - 80s 257ms/step - loss: 0.6702 - accuracy: 0.5985 - val_loss: 0.6055 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.5851 - accuracy: 0.6899
Epoch 00004: val_loss improved from 0.60553 to 0.54122, saving model to best_model_NADAM.h5
313/313 [=====] - 80s 254ms/step - loss: 0.5851 - accuracy: 0.6899 - val_loss: 0.5412 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.5181 - accuracy: 0.7444
Epoch 00005: val_loss improved from 0.54122 to 0.46390, saving model to best_model_NADAM.h5
313/313 [=====] - 80s 255ms/step - loss: 0.5181 - accuracy: 0.7444 - val_loss: 0.4639 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.4558 - accuracy: 0.7839
Epoch 00006: val_loss did not improve from 0.46390
313/313 [=====] - 79s 252ms/step - loss: 0.4558 - accuracy: 0.7839 - val_loss: 0.6407 - val_accuracy: 0.
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.4069 - accuracy: 0.8137
Epoch 00007: val_loss improved from 0.46390 to 0.38349, saving model to best_model_NADAM.h5
313/313 [=====] - 79s 251ms/step - loss: 0.4069 - accuracy: 0.8137 - val_loss: 0.3835 - val_accuracy: 0.
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.3555 - accuracy: 0.8403
Epoch 00008: val_loss improved from 0.38349 to 0.34450, saving model to best_model_NADAM.h5
313/313 [=====] - 78s 250ms/step - loss: 0.3555 - accuracy: 0.8403 - val_loss: 0.3445 - val_accuracy: 0.
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.3301 - accuracy: 0.8531
Epoch 00009: val_loss improved from 0.34450 to 0.29266, saving model to best_model_NADAM.h5
313/313 [=====] - 78s 251ms/step - loss: 0.3301 - accuracy: 0.8531 - val_loss: 0.2927 - val_accuracy: 0.
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.2985 - accuracy: 0.8706
Epoch 00010: val_loss did not improve from 0.29266
313/313 [=====] - 78s 250ms/step - loss: 0.2985 - accuracy: 0.8706 - val_loss: 0.2938 - val_accuracy: 0.
```

```
plt.plot(history_nadam.history['accuracy'])
plt.plot(history_nadam.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model.save_weights('Atharva_Nadam_weights_48.h5')
np.save('Atharva_Nadam_48.npy',history_nadam.history)
```

MODEL 2 - OPTIMIZER : SGD

```
model_sgd = Sequential()
model_sgd.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_sgd.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(MaxPooling2D((2, 2), strides=(2, 2)))
```

```
model_sgd.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgd.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgd.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgd.add(Flatten())

model_sgd.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_sgd.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.SGD(
    learning_rate=0.01, momentum=0.0, nesterov=False)
model_sgd.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_sgd.summary()
```

🔗 Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====		
conv2d_104 (Conv2D)	(None, 48, 48, 256)	7168
<hr/>		
conv2d_105 (Conv2D)	(None, 48, 48, 256)	590080
<hr/>		
max_pooling2d_39 (MaxPooling)	(None, 24, 24, 256)	0
<hr/>		
conv2d_106 (Conv2D)	(None, 24, 24, 128)	295040
<hr/>		
conv2d_107 (Conv2D)	(None, 24, 24, 128)	147584
<hr/>		
max_pooling2d_40 (MaxPooling)	(None, 12, 12, 128)	0
<hr/>		
conv2d_108 (Conv2D)	(None, 12, 12, 64)	73792
<hr/>		
conv2d_109 (Conv2D)	(None, 12, 12, 64)	36928
<hr/>		
conv2d_110 (Conv2D)	(None, 12, 12, 64)	36928
<hr/>		
conv2d_111 (Conv2D)	(None, 12, 12, 64)	36928
<hr/>		
max_pooling2d_41 (MaxPooling)	(None, 6, 6, 64)	0
<hr/>		
flatten_13 (Flatten)	(None, 2304)	0
<hr/>		
dense_26 (Dense)	(None, 128)	295040
<hr/>		
dense_27 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		
<hr/>		

```
from keras.callbacks import ModelCheckpoint


checkpoint = ModelCheckpoint(filepath = 'best_model_SGD.h5',save_best_only = True,verbose=1)

history_sgd = model_sgd.fit(training_set,
                             epochs=10,
                             callbacks=[checkpoint],
                             validation_data=validation_set)
```




```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6796 - accuracy: 0.5662
Epoch 00001: val_loss improved from inf to 0.68032, saving model to best_model_SGD.h5
313/313 [=====] - 79s 251ms/step - loss: 0.6796 - accuracy: 0.5662 - val_loss: 0.6803 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6421 - accuracy: 0.6255
Epoch 00002: val_loss improved from 0.68032 to 0.61192, saving model to best_model_SGD.h5
313/313 [=====] - 79s 251ms/step - loss: 0.6421 - accuracy: 0.6255 - val_loss: 0.6119 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.6091 - accuracy: 0.6681
Epoch 00003: val_loss did not improve from 0.61192
313/313 [=====] - 79s 252ms/step - loss: 0.6091 - accuracy: 0.6681 - val_loss: 0.6181 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.5835 - accuracy: 0.6888
Epoch 00004: val_loss improved from 0.61192 to 0.53752, saving model to best_model_SGD.h5
313/313 [=====] - 78s 250ms/step - loss: 0.5835 - accuracy: 0.6888 - val_loss: 0.5375 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.5612 - accuracy: 0.7095
Epoch 00005: val_loss improved from 0.53752 to 0.51241, saving model to best_model_SGD.h5
313/313 [=====] - 78s 249ms/step - loss: 0.5612 - accuracy: 0.7095 - val_loss: 0.5124 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.5478 - accuracy: 0.7203
Epoch 00006: val_loss did not improve from 0.51241
313/313 [=====] - 78s 249ms/step - loss: 0.5478 - accuracy: 0.7203 - val_loss: 0.5565 - val_accuracy: 0.
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.5281 - accuracy: 0.7372
```

```
plt.plot(history_sgd.history['accuracy'])
plt.plot(history_sgd.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_sgd.save_weights('Atharva_SGD_48.h5')
np.save('Atharva_SGD_48.npy',history_sgd.history)
```

MODEL 3 - OPTIMIZER : SGD + NESTROV

```
model_sgdn = Sequential()
model_sgdn.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_sgdn.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgdn.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_sgdn.add(MaxPooling2D((2, 2), strides=(2,2)))

model_sgdn.add(Flatten())

model_sgdn.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_sgdn.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.SGD(
```

```
learning_rate=0.01, momentum=0.0, nesterov=True)
model_sgdn.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_sgdn.summary()
```

📄

Model: "sequential_14"

Layer (type)	Output Shape	Param #
=====		
conv2d_112 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_113 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_42 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_114 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_115 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_43 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_116 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_117 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_118 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_119 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_44 (MaxPooling)	(None, 6, 6, 64)	0
flatten_14 (Flatten)	(None, 2304)	0
dense_28 (Dense)	(None, 128)	295040
dense_29 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		

```
from keras.callbacks import ModelCheckpoint


checkpoint = ModelCheckpoint(filepath = 'best_model_SGDN.h5',save_best_only = True,verbose=1)

history_sgdn = model_sgdn.fit(training_set,
                               epochs=10,
                               callbacks=[checkpoint],
                               validation_data=validation_set)
```



```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6764 - accuracy: 0.5730
Epoch 00001: val_loss improved from inf to 0.65042, saving model to best_model_SGDN.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6764 - accuracy: 0.5730 - val_loss: 0.6504 - val_accuracy: 0.6311
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6360 - accuracy: 0.6311

plt.plot(history_sgdn.history['accuracy'])
plt.plot(history_sgdn.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.6260 - accuracy: 0.6260
Epoch 00009: val_loss improved from 0.65042 to 0.62602, saving model to best_model_SGDN.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6260 - accuracy: 0.6260 - val_loss: 0.6260 - val_accuracy: 0.6311
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.6260 - accuracy: 0.6311
Epoch 00010: val_loss did not improve, saving model to best_model_SGDN.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6260 - accuracy: 0.6311 - val_loss: 0.6260 - val_accuracy: 0.6311

model_sgdn.save_weights('Atharva_SGDN_48.h5')
np.save('Atharva_SGDN_48.npy',history_sgdn.history)
```

MODEL 4 - OPTIMIZER : RMS

```
model_rms = Sequential()
model_rms.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_rms.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(MaxPooling2D((2, 2), strides=(2,2)))

model_rms.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(MaxPooling2D((2, 2), strides=(2,2)))

model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms.add(MaxPooling2D((2, 2), strides=(2,2)))

model_rms.add(Flatten())

model_rms.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_rms.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.RMSprop(
    learning_rate=0.001,
    rho=0.9,
    momentum=0.0,
    epsilon=1e-07,
    centered=False,)
model_rms.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_rms.summary()
```



Model: "sequential_15"

Layer (type)	Output Shape	Param #
=====		
conv2d_120 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_121 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_45 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_122 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_123 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_46 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_124 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_125 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_126 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_127 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_47 (MaxPooling)	(None, 6, 6, 64)	0


```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath = 'best_model_RMS.h5',save_best_only = True,verbose=1)
```

```
history_rms = model_rms.fit(training_set,
                             epochs=10,
                             callbacks=[checkpoint],
                             validation_data=validation_set)
```

```
🔗 Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 1.1355 - accuracy: 0.5162
Epoch 00001: val_loss improved from inf to 0.69124, saving model to best_model_RMS.h5
313/313 [=====] - 78s 248ms/step - loss: 1.1355 - accuracy: 0.5162 - val_loss: 0.6912 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6784 - accuracy: 0.5974
Epoch 00002: val_loss improved from 0.69124 to 0.59583, saving model to best_model_RMS.h5
313/313 [=====] - 77s 246ms/step - loss: 0.6784 - accuracy: 0.5974 - val_loss: 0.5958 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.5916 - accuracy: 0.6934
Epoch 00003: val_loss improved from 0.59583 to 0.51424, saving model to best_model_RMS.h5
313/313 [=====] - 78s 248ms/step - loss: 0.5916 - accuracy: 0.6934 - val_loss: 0.5142 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.5089 - accuracy: 0.7548
Epoch 00004: val_loss improved from 0.51424 to 0.47275, saving model to best_model_RMS.h5
313/313 [=====] - 78s 248ms/step - loss: 0.5089 - accuracy: 0.7548 - val_loss: 0.4728 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.4463 - accuracy: 0.7972
Epoch 00005: val_loss improved from 0.47275 to 0.37714, saving model to best_model_RMS.h5
313/313 [=====] - 77s 246ms/step - loss: 0.4463 - accuracy: 0.7972 - val_loss: 0.3771 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.4009 - accuracy: 0.8160
Epoch 00006: val_loss improved from 0.37714 to 0.34803, saving model to best_model_RMS.h5
313/313 [=====] - 77s 246ms/step - loss: 0.4009 - accuracy: 0.8160 - val_loss: 0.3480 - val_accuracy: 0.
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.3708 - accuracy: 0.8323
Epoch 00007: val_loss improved from 0.34803 to 0.31253, saving model to best_model_RMS.h5
313/313 [=====] - 77s 246ms/step - loss: 0.3708 - accuracy: 0.8323 - val_loss: 0.3125 - val_accuracy: 0.
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.3437 - accuracy: 0.8490
Epoch 00008: val_loss improved from 0.31253 to 0.31074, saving model to best_model_RMS.h5
313/313 [=====] - 77s 245ms/step - loss: 0.3437 - accuracy: 0.8490 - val_loss: 0.3107 - val_accuracy: 0.
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.3290 - accuracy: 0.8562
Epoch 00009: val_loss did not improve from 0.31074
313/313 [=====] - 77s 245ms/step - loss: 0.3290 - accuracy: 0.8562 - val_loss: 0.3168 - val_accuracy: 0.
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.3126 - accuracy: 0.8663
Epoch 00010: val_loss improved from 0.31074 to 0.28536, saving model to best_model_RMS.h5
313/313 [=====] - 77s 245ms/step - loss: 0.3126 - accuracy: 0.8663 - val_loss: 0.2854 - val_accuracy: 0.
```

```
plt.plot(history_rms.history['accuracy'])
plt.plot(history_rms.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_rms.save_weights('Atharva_RMS_4850.h5')
np.save('Atharva_RMS_4850.npy',history_rms.history)
```

MODEL 5 - OPTIMIZER : ADAM

```
model_adam = Sequential()
model_adam.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adam.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adam.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adam.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adam.add(Flatten())

model_adam.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adam.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Adam(
    learning_rate=0.001,
    beta_1=0.9,
    beta_2=0.999,
    epsilon=1e-07,
    amsgrad=False,)
model_adam.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_adam.summary()
```



Model: "sequential_16"

Layer (type)	Output Shape	Param #
=====		
conv2d_128 (Conv2D)	(None, 48, 48, 256)	7168

conv2d_129 (Conv2D)	(None, 48, 48, 256)	590080

max_pooling2d_48 (MaxPooling)	(None, 24, 24, 256)	0


```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath = 'best_model_ADAM.h5',save_best_only = True,verbose=1)
```

```
history_adam = model_adam.fit(training_set,
                               epochs=10,
                               callbacks=[checkpoint],
                               validation_data=validation_set)
```

```
➤ Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6936 - accuracy: 0.5601
Epoch 00001: val_loss improved from inf to 0.65515, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 247ms/step - loss: 0.6936 - accuracy: 0.5601 - val_loss: 0.6551 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6354 - accuracy: 0.6440
Epoch 00002: val_loss improved from 0.65515 to 0.58563, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.6354 - accuracy: 0.6440 - val_loss: 0.5856 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.5680 - accuracy: 0.7085
Epoch 00003: val_loss improved from 0.58563 to 0.50897, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.5680 - accuracy: 0.7085 - val_loss: 0.5090 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.4974 - accuracy: 0.7588
Epoch 00004: val_loss improved from 0.50897 to 0.44578, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.4974 - accuracy: 0.7588 - val_loss: 0.4458 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.4427 - accuracy: 0.7914
Epoch 00005: val_loss improved from 0.44578 to 0.39582, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 246ms/step - loss: 0.4427 - accuracy: 0.7914 - val_loss: 0.3958 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.3863 - accuracy: 0.8246
Epoch 00006: val_loss improved from 0.39582 to 0.34291, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 245ms/step - loss: 0.3863 - accuracy: 0.8246 - val_loss: 0.3429 - val_accuracy: 0.
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.3628 - accuracy: 0.8396
Epoch 00007: val_loss did not improve from 0.34291
313/313 [=====] - 77s 245ms/step - loss: 0.3628 - accuracy: 0.8396 - val_loss: 0.3883 - val_accuracy: 0.
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.3288 - accuracy: 0.8512
Epoch 00008: val_loss improved from 0.34291 to 0.29846, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 246ms/step - loss: 0.3288 - accuracy: 0.8512 - val_loss: 0.2985 - val_accuracy: 0.
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.3052 - accuracy: 0.8667
Epoch 00009: val_loss improved from 0.29846 to 0.28250, saving model to best_model_ADAM.h5
313/313 [=====] - 77s 247ms/step - loss: 0.3052 - accuracy: 0.8667 - val_loss: 0.2825 - val_accuracy: 0.
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.2934 - accuracy: 0.8735
Epoch 00010: val_loss improved from 0.28250 to 0.26183, saving model to best_model_ADAM.h5
313/313 [=====] - 78s 248ms/step - loss: 0.2934 - accuracy: 0.8735 - val_loss: 0.2618 - val_accuracy: 0.
```

```
plt.plot(history_adam.history['accuracy'])
plt.plot(history_adam.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adam.save_weights('Atharva_ADAM_48.h5')
np.save('Atharva_ADAM_48.npy', history_adam.history)
```

MODEL 6 - OPTIMIZER : ADADelta

```
model_adad = Sequential()
model_adad.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adad.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adad.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adad.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adad.add(Flatten())

model_adad.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adad.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Adadelta(
    learning_rate=0.001, rho=0.95, epsilon=1e-07, name="Adadelta", )
model_adad.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_adad.summary()
```

```

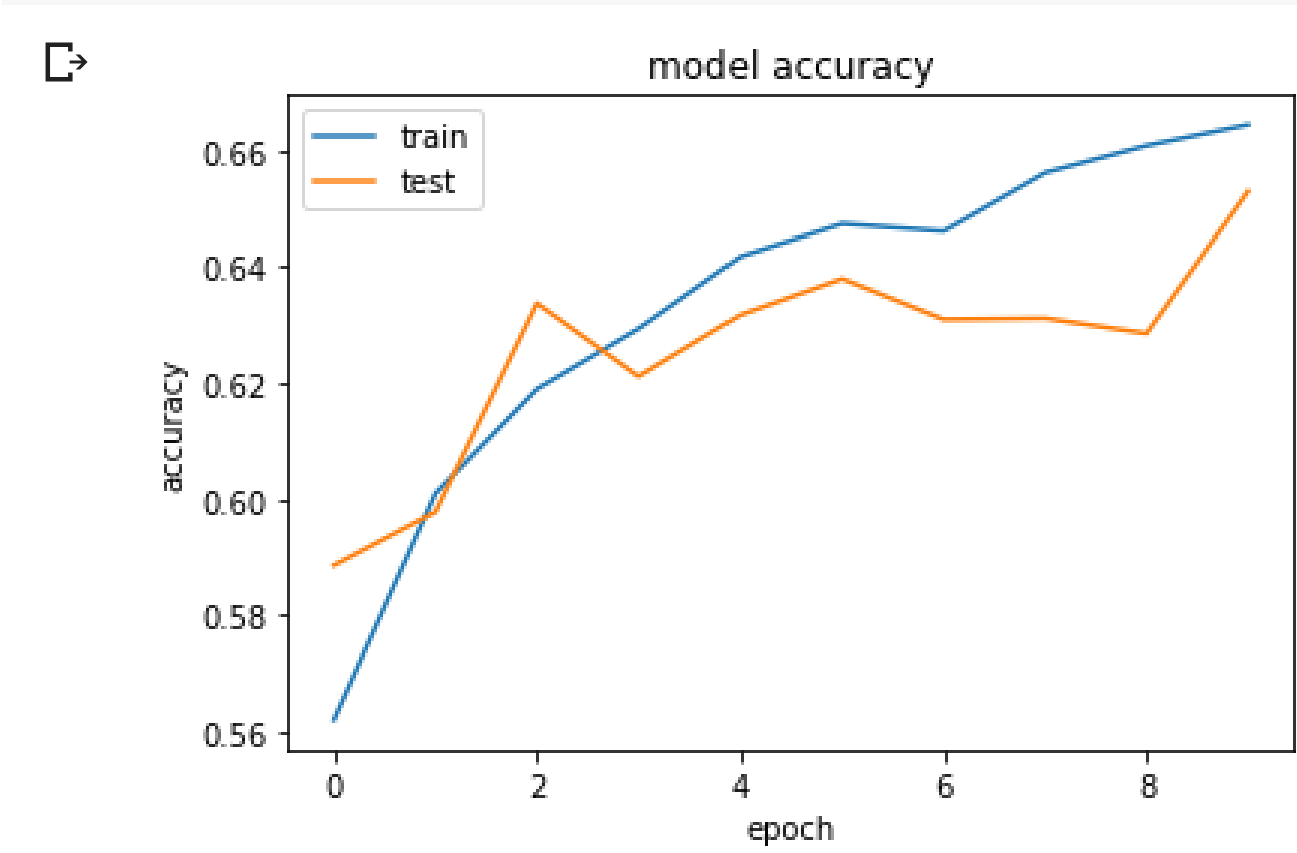
Model: "sequential_17"
Layer (type)                Output Shape                Param #
=====
conv2d_136 (Conv2D)          (None, 48, 48, 256)        7168
conv2d_137 (Conv2D)          (None, 48, 48, 256)        590080
max_pooling2d_51 (MaxPooling (None, 24, 24, 256)        0
conv2d_138 (Conv2D)          (None, 24, 24, 128)        295040
conv2d_139 (Conv2D)          (None, 24, 24, 128)        147584
max_pooling2d_52 (MaxPooling (None, 12, 12, 128)        0
conv2d_140 (Conv2D)          (None, 12, 12, 64)         73792
conv2d_141 (Conv2D)          (None, 12, 12, 64)         36928
conv2d_142 (Conv2D)          (None, 12, 12, 64)         36928
conv2d_143 (Conv2D)          (None, 12, 12, 64)         36928
max_pooling2d_53 (MaxPooling (None, 6, 6, 64)           0
flatten_17 (Flatten)         (None, 2304)                0
dense_34 (Dense)             (None, 128)                 295040
dense_35 (Dense)             (None, 1)                   129
=====
Total params: 1,519,617
Trainable params: 1,519,617
Non-trainable params: 0

```

[illegible]

```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6824 - accuracy: 0.5621
Epoch 00001: val_loss improved from inf to 0.67433, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 251ms/step - loss: 0.6824 - accuracy: 0.5621 - val_loss: 0.6743 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6702 - accuracy: 0.6010
Epoch 00002: val_loss improved from 0.67433 to 0.66563, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6702 - accuracy: 0.6010 - val_loss: 0.6656 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.6615 - accuracy: 0.6190
Epoch 00003: val_loss improved from 0.66563 to 0.65408, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6615 - accuracy: 0.6190 - val_loss: 0.6541 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.6528 - accuracy: 0.6294
Epoch 00004: val_loss improved from 0.65408 to 0.65017, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6528 - accuracy: 0.6294 - val_loss: 0.6502 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.6444 - accuracy: 0.6418
Epoch 00005: val_loss improved from 0.65017 to 0.64251, saving model to best_model_ADADELTA.h5
313/313 [=====] - 80s 255ms/step - loss: 0.6444 - accuracy: 0.6418 - val_loss: 0.6425 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.6373 - accuracy: 0.6474
Epoch 00006: val_loss improved from 0.64251 to 0.63577, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6373 - accuracy: 0.6474 - val_loss: 0.6358 - val_accuracy: 0.
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.6325 - accuracy: 0.6461
Epoch 00007: val_loss improved from 0.63577 to 0.63574, saving model to best_model_ADADELTA.h5
313/313 [=====] - 80s 254ms/step - loss: 0.6325 - accuracy: 0.6461 - val_loss: 0.6357 - val_accuracy: 0.
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.6269 - accuracy: 0.6562
Epoch 00008: val_loss improved from 0.63574 to 0.63433, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 253ms/step - loss: 0.6269 - accuracy: 0.6562 - val_loss: 0.6343 - val_accuracy: 0.
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.6213 - accuracy: 0.6608
Epoch 00009: val_loss did not improve from 0.63433
313/313 [=====] - 79s 253ms/step - loss: 0.6213 - accuracy: 0.6608 - val_loss: 0.6356 - val_accuracy: 0.
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.6180 - accuracy: 0.6644
Epoch 00010: val_loss improved from 0.63433 to 0.61984, saving model to best_model_ADADELTA.h5
313/313 [=====] - 79s 252ms/step - loss: 0.6180 - accuracy: 0.6644 - val_loss: 0.6198 - val_accuracy: 0.
```

```
plt.plot(history_adad.history['accuracy'])
plt.plot(history_adad.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adad.save_weights('Atharva_ADADELTA_48.h5')
np.save('Atharva_ADADELTA_48.npy',history_adad.history)
```

MODEL 7 - OPTIMIZER : ADAGrad

```
model_adag = Sequential()
model_adag.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adag.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_adag.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(MaxPooling2D((2, 2), strides=(2,2)))
```

```
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
```



```
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adag.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adag.add(Flatten())

model_adag.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adag.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Adagrad(
    learning_rate=0.001,
    initial_accumulator_value=0.1,
    epsilon=1e-07,)
model_adag.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_adag.summary()
```

🔗 Model: "sequential_18"

Layer (type)	Output Shape	Param #
=====		
conv2d_144 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_145 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_54 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_146 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_147 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_55 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_148 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_149 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_150 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_151 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_56 (MaxPooling)	(None, 6, 6, 64)	0
flatten_18 (Flatten)	(None, 2304)	0
dense_36 (Dense)	(None, 128)	295040
dense_37 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		

```
from keras.callbacks import ModelCheckpoint


checkpoint = ModelCheckpoint(filepath = 'best_model_ADAGRAD.h5',save_best_only = True,verbose=1)

history_adag = model_adag.fit(training_set,
                               epochs=10,
                               callbacks=[checkpoint],
                               validation_data=validation_set)
```




```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6747 - accuracy: 0.5795
Epoch 00001: val_loss improved from inf to 0.64739, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 78s 250ms/step - loss: 0.6747 - accuracy: 0.5795 - val_loss: 0.6474 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6404 - accuracy: 0.6321
Epoch 00002: val_loss did not improve from 0.64739
313/313 [=====] - 78s 250ms/step - loss: 0.6404 - accuracy: 0.6321 - val_loss: 0.7419 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.6191 - accuracy: 0.6536
Epoch 00003: val_loss improved from 0.64739 to 0.58838, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 78s 248ms/step - loss: 0.6191 - accuracy: 0.6536 - val_loss: 0.5884 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.6012 - accuracy: 0.6745
Epoch 00004: val_loss did not improve from 0.58838
313/313 [=====] - 78s 250ms/step - loss: 0.6012 - accuracy: 0.6745 - val_loss: 0.6177 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.5848 - accuracy: 0.6920
Epoch 00005: val_loss improved from 0.58838 to 0.57214, saving model to best_model_ADAGRAD.h5
313/313 [=====] - 79s 252ms/step - loss: 0.5848 - accuracy: 0.6920 - val_loss: 0.5721 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.5678 - accuracy: 0.7046
```

```
plt.plot(history_adag.history['accuracy'])
plt.plot(history_adag.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_adag.save_weights('Atharva_ADAG_48.h5')
np.save('Atharva_ADAG_48.npy',history_adag.history)
```

MODEL 8 - OPTIMIZER : ADAMax

```
model_adamax = Sequential()
model_adamax.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_adamax.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adamax.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_adamax.add(MaxPooling2D((2, 2), strides=(2,2)))

model_adamax.add(Flatten())

model_adamax.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))

model_adamax.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Adamax(
    learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07,)
model_adamax.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_adamax.summary()
```

Model: "sequential_19"

Layer (type)	Output Shape	Param #
=====		
conv2d_152 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_153 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_57 (MaxPooling)	(None, 24, 24, 256)	0
conv2d_154 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_155 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_58 (MaxPooling)	(None, 12, 12, 128)	0
conv2d_156 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_157 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_158 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_159 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_59 (MaxPooling)	(None, 6, 6, 64)	0
flatten_19 (Flatten)	(None, 2304)	0
dense_38 (Dense)	(None, 128)	295040
dense_39 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		

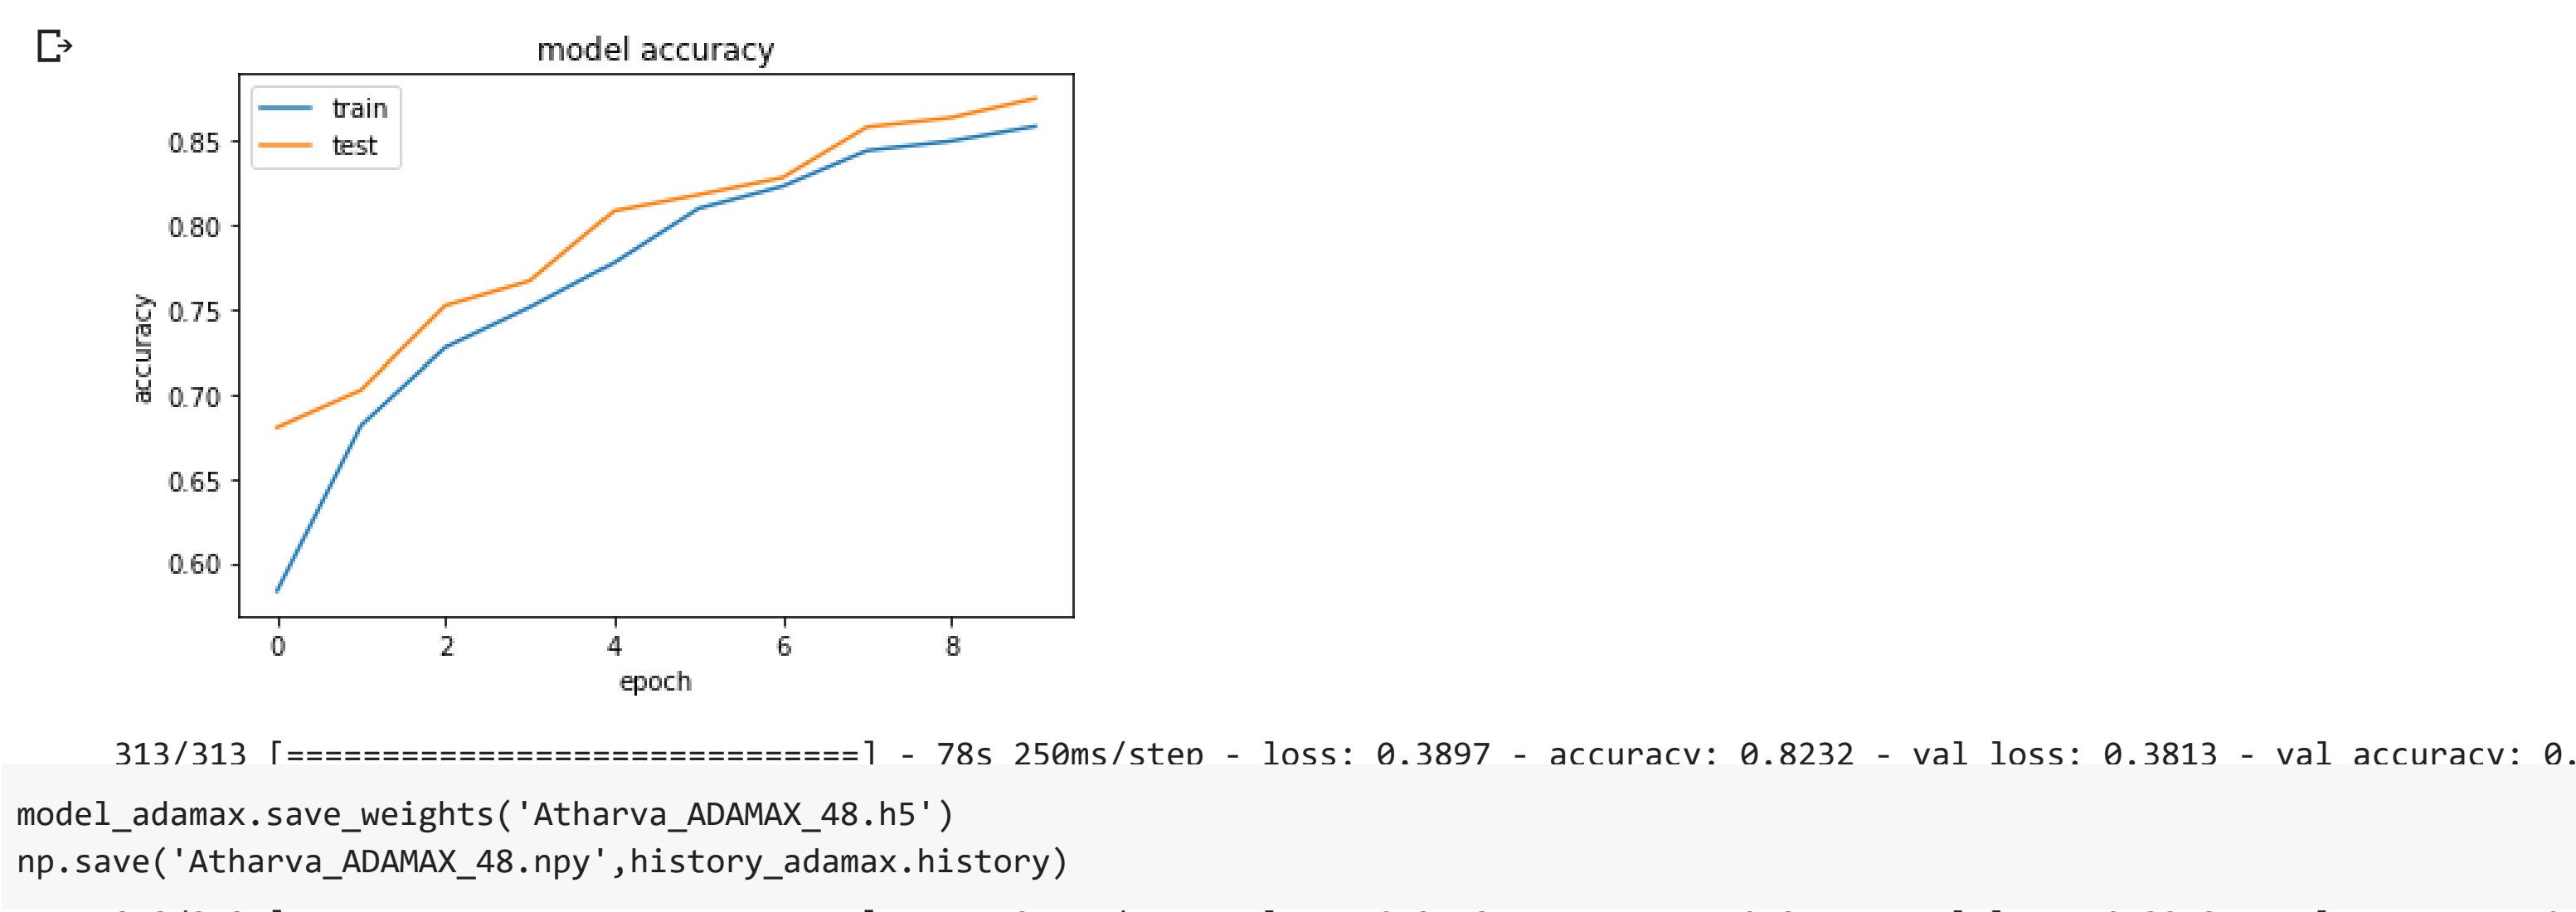
```
from keras.callbacks import ModelCheckpoint

checkpoint = ModelCheckpoint(filepath = 'best_model_ADAMAX.h5',save_best_only = True,verbose=1)

history_adamax = model_adamax.fit(training_set,
                                   epochs=10,
                                   callbacks=[checkpoint],
                                   validation_data=validation_set)
```



```
Epoch 1/10
plt.plot(history_adamax.history[ 'accuracy' ])
plt.plot(history_adamax.history[ 'val_accuracy' ])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



Testing the Model against custom test data

We will validate the test accuracy using the model generated using the NADAM optimizer as it has the highest degree of performance

```
Epoch 00010: val loss improved from 0.31083 to 0.28348 - saving model to best_model_ADAMAX.h5
!mkdir augmentedtest
```

```
test_datagen = ImageDataGenerator(rescale = 1.0/255.0,
                                  shear_range = 0.2,
                                  zoom_range = 0.2,
                                  horizontal_flip = True,
                                  )

test_set = test_datagen.flow_from_directory('./test/',
                                           target_size = (48, 48),
                                           batch_size = 1,
                                           class_mode = 'binary',
                                           save_to_dir='./augmentedtest',
                                           save_prefix='aug',
                                           save_format='jpg')

Found 40 images belonging to 1 classes.
```

```
import cv2
predict_files = glob.glob("/content/augmentedtest/*.jpg")
predictor, image_id = [], []
for file in predict_files:
    image = cv2.imread(file)
    img = cv2.resize(image, (48,48))
    img = np.expand_dims(img, 0)

    prediction = model.predict(img)

    predict = str(model.predict_classes(img)[0][0])

    predictor.extend(list(predict))
```

```
predictions =[]
c = 0
w = 0
classified = ['cat','dog']
for i in range(len(predictor)):
    if predictor[i] == 0:
        predictions.append(classified[0])
    else:
        predictions.append(classified[1])
```

```
predict_files = glob.glob("/content/augmentedtest*.jpg*")
```

```
for file in predict_files:
```

```
data.append(os.path.basename(file)[:3])
```

```
data = pd.DataFrame(data, columns=['FileName'])
```

```
data['Predictions'] = predictions
```


Last half images are of dogs - so by using data.tail we can see the predicted results

```
data.tail(10)
```

	FileName
150	cat
151	dog
152	dog
153	dog
154	cat
155	dog
156	dog
157	cat
158	dog
159	cat

First half images are of cats - so by using data.head we can see the predicted results

```
data.head(10)
```

	FileName
0	dog
1	cat
2	cat
3	dog
4	dog
5	cat
6	cat
7	cat
8	cat
9	dog

Regularized Model 1 - Optimizer : NADAM

Since Nadam and RMSProp performed the best in unregularised models, we will only try regularisation on them.

```
model_nadam_r = Sequential()
model_nadam_r.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_nadam_r.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_nadam_r.add(MaxPooling2D((2, 2), strides=(2,2)))
model_nadam_r.add(Dropout(0.1))
model_nadam_r.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_nadam_r.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_nadam_r.add(MaxPooling2D((2, 2), strides=(2,2)))
model_nadam_r.add(Dropout(0.1))
model_nadam_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_nadam_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_nadam_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_nadam_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
```

```
model_nadam_r.add(MaxPooling2D((2, 2), strides=(2,2)))
model_nadam_r.add(Dropout(0.2))
model_nadam_r.add(Flatten())
model_nadam_r.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
model_nadam_r.add(Dense(1, activation='sigmoid'))

opt = tf.keras.optimizers.Nadam(
    learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07)
model_nadam_r.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

model_nadam_r.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 48, 48, 256)	7168
conv2d_9 (Conv2D)	(None, 48, 48, 256)	590080
max_pooling2d_3 (MaxPooling2D)	(None, 24, 24, 256)	0
dropout (Dropout)	(None, 24, 24, 256)	0
conv2d_10 (Conv2D)	(None, 24, 24, 128)	295040
conv2d_11 (Conv2D)	(None, 24, 24, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 12, 12, 128)	0
dropout_1 (Dropout)	(None, 12, 12, 128)	0
conv2d_12 (Conv2D)	(None, 12, 12, 64)	73792
conv2d_13 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_14 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_15 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_5 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_2 (Dropout)	(None, 6, 6, 64)	0
flatten_1 (Flatten)	(None, 2304)	0
dense_2 (Dense)	(None, 128)	295040
dense_3 (Dense)	(None, 1)	129
=====		
Total params: 1,519,617		
Trainable params: 1,519,617		
Non-trainable params: 0		
=====		

```
from keras.callbacks import ModelCheckpoint


checkpoint = ModelCheckpoint(filepath = 'best_model_NADAM-Regularized.h5',save_best_only = True,verbose=1)

history_nadam_r = model.fit(training_set,
                             epochs=10,
                             callbacks=[checkpoint],
                             validation_data=validation_set)
```



```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.6989 - accuracy: 0.5597
Epoch 00001: val_loss improved from inf to 0.65994, saving model to best_model_NADAM-Regularized.h5
313/313 [=====] - 71s 228ms/step - loss: 0.6989 - accuracy: 0.5597 - val_loss: 0.6599 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.6613 - accuracy: 0.6120
Epoch 00002: val_loss improved from 0.65994 to 0.61558, saving model to best_model_NADAM-Regularized.h5
313/313 [=====] - 71s 227ms/step - loss: 0.6613 - accuracy: 0.6120 - val_loss: 0.6156 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.5822 - accuracy: 0.6925
Epoch 00003: val_loss improved from 0.61558 to 0.51408, saving model to best_model_NADAM-Regularized.h5
313/313 [=====] - 71s 227ms/step - loss: 0.5822 - accuracy: 0.6925 - val_loss: 0.5141 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.5053 - accuracy: 0.7531
Epoch 00004: val_loss improved from 0.51408 to 0.45279, saving model to best_model_NADAM-Regularized.h5
313/313 [=====] - 71s 227ms/step - loss: 0.5053 - accuracy: 0.7531 - val_loss: 0.4528 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.4461 - accuracy: 0.7915
Epoch 00005: val loss did not improve from 0.45279
```

```
plt.plot(history_nadam_r.history['accuracy'])
plt.plot(history_nadam_r.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_nadam_r.save_weights('Atharva_NADAM-Regularized_48.h5')
np.save('Atharva_NADAM-Regularized_48.npy',history_nadam_r.history)
```

Regularized Model 2 - Optimizer : RMSProp

Since Nadam and RMSProp performed the best in unregularised models, we will only try regularisation on them.

```
model_rms_r = Sequential()
model_rms_r.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(48, 48, 3)))
model_rms_r.add(Conv2D(256, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(MaxPooling2D((2, 2), strides=(2,2)))
model_rms_r.add(Dropout(0.1))
model_rms_r.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(MaxPooling2D((2, 2), strides=(2,2)))
model_rms_r.add(Dropout(0.1))
model_rms_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same'))
model_rms_r.add(MaxPooling2D((2, 2), strides=(2,2)))
model_rms_r.add(Dropout(0.2))
model_rms_r.add(Flatten())
model_rms_r.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
model_rms_r.add(Dense(1, activation='sigmoid'))
```


```
opt = tf.keras.optimizers.RMSprop(
    learning_rate=0.001,
    rho=0.9,
    momentum=0.0,
    epsilon=1e-07,
    centered=False,)
model_rms_r.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```



```
checkpoint = ModelCheckpoint(filepath = 'best_model_RMS-Regularized.h5',save_best_only = True,verbose=1)
history_rms_r = model.fit(training_set,
                           epochs=10,
                           callbacks=[checkpoint],
                           validation_data=validation_set)
```

```
Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.2744 - accuracy: 0.8832
Epoch 00001: val_loss improved from inf to 0.25109, saving model to best_model_RMS-Regularized.h5
313/313 [=====] - 71s 228ms/step - loss: 0.2744 - accuracy: 0.8832 - val_loss: 0.2511 - val_accuracy: 0.
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.2594 - accuracy: 0.8881
Epoch 00002: val_loss did not improve from 0.25109
313/313 [=====] - 71s 226ms/step - loss: 0.2594 - accuracy: 0.8881 - val_loss: 0.3113 - val_accuracy: 0.
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.2435 - accuracy: 0.8969
Epoch 00003: val_loss did not improve from 0.25109
313/313 [=====] - 71s 225ms/step - loss: 0.2435 - accuracy: 0.8969 - val_loss: 0.2599 - val_accuracy: 0.
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.2358 - accuracy: 0.8989
Epoch 00004: val_loss improved from 0.25109 to 0.23949, saving model to best_model_RMS-Regularized.h5
313/313 [=====] - 71s 226ms/step - loss: 0.2358 - accuracy: 0.8989 - val_loss: 0.2395 - val_accuracy: 0.
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.2252 - accuracy: 0.9049
Epoch 00005: val_loss did not improve from 0.23949
313/313 [=====] - 71s 226ms/step - loss: 0.2252 - accuracy: 0.9049 - val_loss: 0.2398 - val_accuracy: 0.
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.2121 - accuracy: 0.9086
Epoch 00006: val_loss improved from 0.23949 to 0.23455, saving model to best_model_RMS-Regularized.h5
313/313 [=====] - 71s 227ms/step - loss: 0.2121 - accuracy: 0.9086 - val_loss: 0.2345 - val_accuracy: 0.
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.2076 - accuracy: 0.9115
Epoch 00007: val_loss did not improve from 0.23455
313/313 [=====] - 71s 226ms/step - loss: 0.2076 - accuracy: 0.9115 - val_loss: 0.2557 - val_accuracy: 0.
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.2012 - accuracy: 0.9163
Epoch 00008: val_loss did not improve from 0.23455
313/313 [=====] - 71s 225ms/step - loss: 0.2012 - accuracy: 0.9163 - val_loss: 0.2431 - val_accuracy: 0.
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.1953 - accuracy: 0.9183
Epoch 00009: val_loss did not improve from 0.23455
313/313 [=====] - 70s 225ms/step - loss: 0.1953 - accuracy: 0.9183 - val_loss: 0.2637 - val_accuracy: 0.
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.1875 - accuracy: 0.9216
Epoch 00010: val_loss improved from 0.23455 to 0.23139, saving model to best_model_RMS-Regularized.h5
313/313 [=====] - 71s 227ms/step - loss: 0.1875 - accuracy: 0.9216 - val_loss: 0.2314 - val_accuracy: 0.
```

```
plt.plot(history_rms_r.history['accuracy'])
plt.plot(history_rms_r.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
model_rms_r.save_weights('Atharva_RMS-Regularized_48.h5')
np.save('Atharva_RMS-Regularized_48.npy',history_rms_r.history)
```

CONCLUSION

Thus, we successfully implemented Image Classification using CatsvsDogs Dataset